

NETWORKING!

A COMPUTER NETWORKING ZINE!

BY JULIA EVANS!

like this?
you can print more!
for free!
<http://jvns.ca/zines>

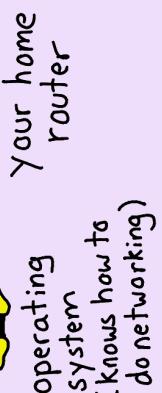


cast of characters

in your house



your laptop
(that you use
to look at **cats**)
cool.py
your
program
(knows how to
do networking)



computes you'll talk to

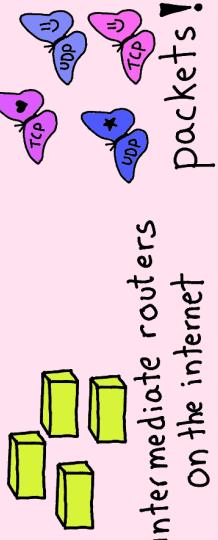


jvns.ca
server
(has **cat**
picture)
DNS server
(Knows which
server hosts
jvns.ca)



the **cat** picture
we're downloading

in the middle

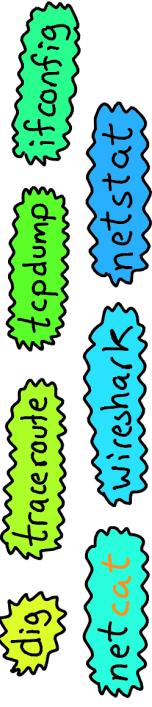


intermediate routers
on the internet
packets!

thanks
for reading

If you want to know more about networking:

→ make network requests! play with



→ beej's guide to network programming is

a useful + funny guide to the socket
API on Unix systems.

→ beej.us/guide/bagnet

→ High Performance Browser Networking
is a **fantastic** ★ and practical guide
to what you need to know about networking
to make fast websites.

You can read it for free at:

→ hpbn.co

Thanks to Kamal Marhubi, Chris Kanich, and
Ada Munroe for reviewing this!

Cover art by the amazing Liz Baillie

Wireshark

What's this?!

Wireshark is an **amazing** tool for packet analysis. Here's an exercise to learn it! Run this:

```
sudo tcpdump port 80 -w http.pcap
```

While that's running, open metafilter.com in your browser. Then press Ctrl+C to stop tcpdump. Now we have a pcap!

Open http.pcap with Wireshark.

Some questions you can try to answer:

- ① What HTTP headers did your browser send to metafilter.com?
(hint: search frame contains "GET")

- ② How many packets were exchanged with metafilter.com's server?
(hint: search 'ip.dst == 54.1.2.3') put the IP from "ping metafilter.com" here

Wireshark makes it easy to look at:

- IP addresses and ports
- **SYNs** and **ACKs** for **TCP** traffic
- exactly what's happening with **DNS** requests
- and so much more. It's a great way to poke around and learn.

hi! I'm Julia

twitter: @b0rk

blog: <http://jvns.ca>

I put a picture of a **cat** on the internet here:

jvns.ca/cat.png ★ (go look!)

In this zine we'll learn everything (mostly) that needs to happen to get that **cat** picture from my server to your laptop.

My goal is to help get you from

"I've heard about some of these HTTP/DNS/TCP things but I don't understand how they work exactly or how they all fit together"

me after I'd been working as a web developer for a year

to...

"there's a networking problem! I totally know where to start!"

me now

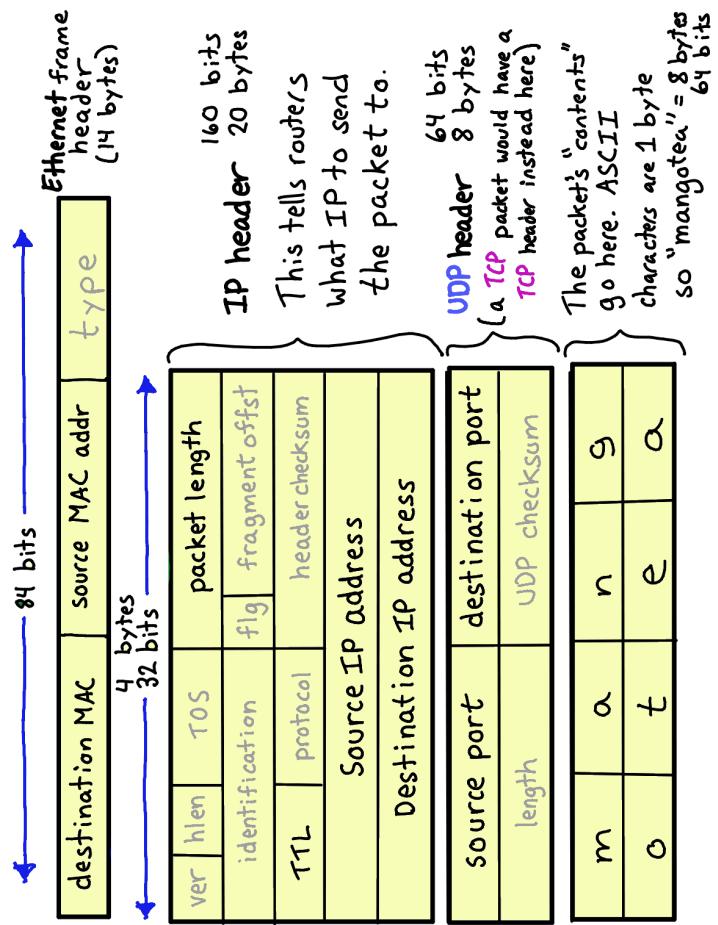
our star: the packet

All data is sent over the internet in **packets**. A packet is a series of bits (0s and 1s...) and it's split into sections (or "headers")

Here's what a **UDP** packet that says "mangotea" looks like. It's 50 bytes in all! (400 bits)

Julia I don't understand this diagram

We are going to work on explaining it!



SSL / TLS

(**TLS**: newer version of **SSL**)

When you send a packet on the internet, LOTS of people can potentially read it.

CAFFÉ
that person is sending email with pie recipes HMMH.

unencrypted wifi

SSL encrypts your packets:

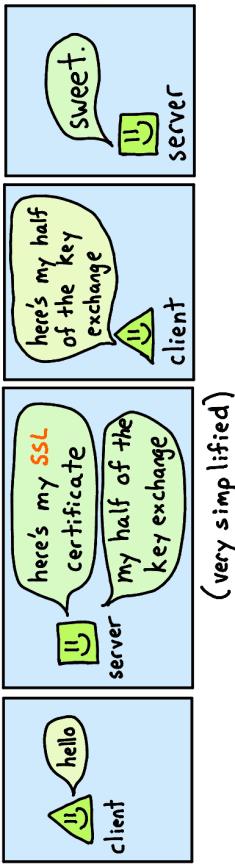
Old packet IP address+port **new packet**

to: 9.9.32.94:443 stay the same to: 9.9.32.94:443 is the usual **SSL** part

from: 31.99.1.2:999 from: 31.99.1.2:999

here is my secret lemon pie recipe \Rightarrow $x\ 8;\ fae\ 94\ aex\ jibq\ 3; 8\ b;" 5jk\$ nobody's gonna know the secret pie recipe NOW!

What happens when you go to <https://jvns.ca>:



Once the client and server agree on a key for the session, they can encrypt all the communication they want.

To see the certificate for jvns.ca, run:

```
$ openssl s_client -connect jvns.ca:443 -servername jvns.ca
```

TLS is really complicated. You can use a tool like **SSL Labs** to check the security of your site.

so "mangotea" = 8 bytes
characters are 1 byte
so "mangotea" = 8 bytes
64 bits

Notation time!

10.0.0.0/8 132.5.23.0/24

People describe groups of IP addresses using CIDR notation.

Example CIDRs

CIDR	range of IPs	
10.0.0.0/8	10.*.*.*	
10.9.0.0/16	10.9.*.*	and 172.16.0.0/12 are reserved for local networking.
10.9.8.0/24	10.9.8.*	

Important examples

In CIDR notation, a /n gives you 2^{32-n} IP addresses. So a /24 is $2^8 = 256$ IPs.

It's important to represent groups of IP addresses efficiently because routers have LOTS TO DO.

is 192.168.3.2 in the subnet 192.168.0.0/16? I can do some really fast bit arithmetic and find out !

10.9.0.0 is this in binary:

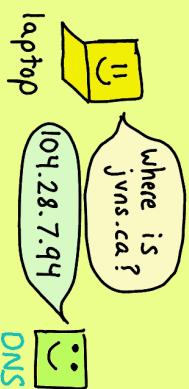
00001010 00001001 00000000 00000000
first 24 bits

10.9.0.0/24 is all the IP addresses which have the same first 24 bits as 10.9.0.0!

steps to get a cat picture from jvns.ca/cat.png

When you download an image, there are a LOT of networking moving pieces. Here are the basic steps we'll explain in the next few pages.

① get the IP address for jvns.ca



③ open a TCP connection to 104.28.7.94 port 80



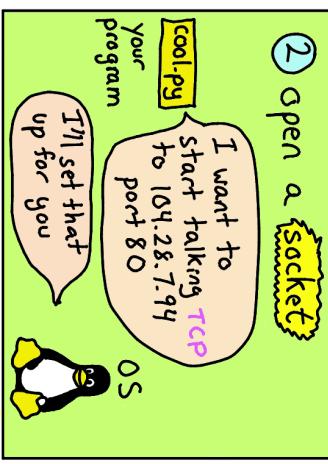
this is a "TCP handshake"
we'll explain it in the TCP section

⑤ get a cat back

HTTP/1.1 200 OK
Content-Type: image/png
Content-Length: 123456
<PNG Bytes>

server with cats

② open a socket



④ request a cat



⑥ clean up

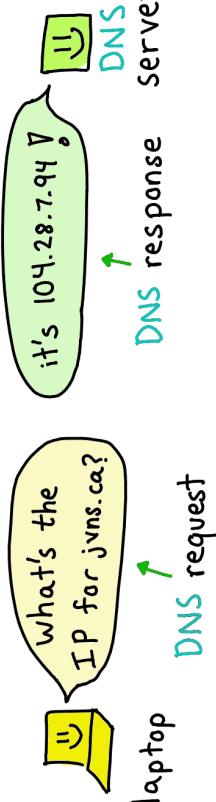
- close the connection maybe
- put the bytes for the PNG in a file maybe
- look at cats definitely



DNS

★ ★ Step ① : get the IP address for jvns.ca ★ ★

All networking happens by sending packets. To send a packet to a server on the internet, you need an **IP address** like 104.28.7.94 jvns.ca and google.com are domain names. **DNS** (the "Domain Name System") is the protocol we use to get the IP address for a domain name.



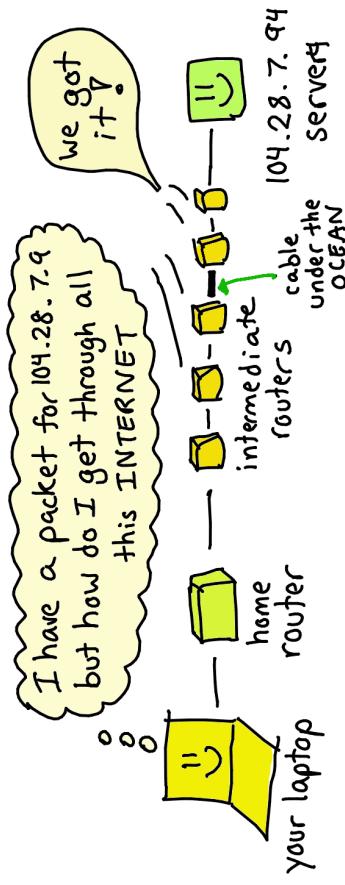
The **DNS** request + response are both usually **UDP** packets.



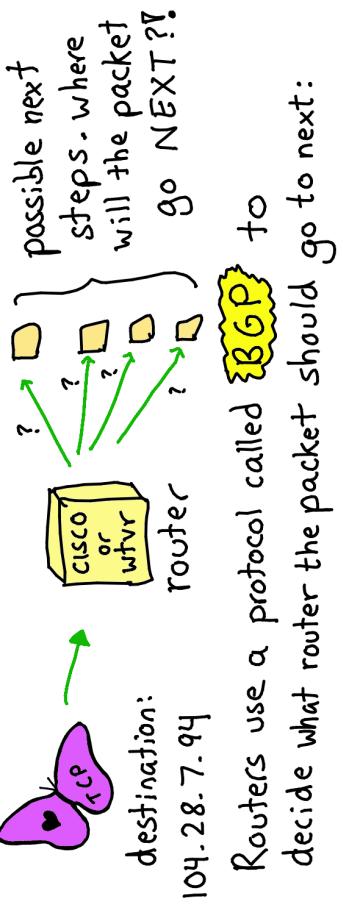
Your system's default **DNS** server is often configured in /etc/resolv.conf.

8.8.8 is Google's **DNS** server, and lots of people use it. It's a great choice!

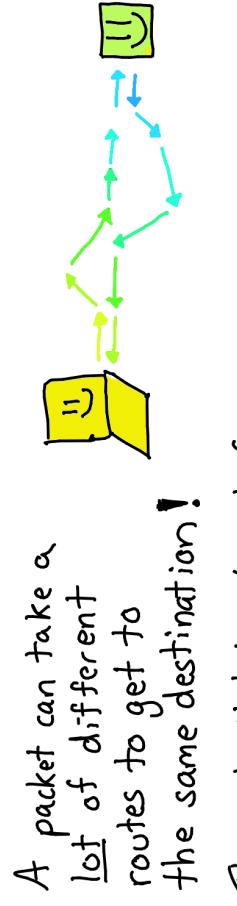
How packets get sent across the ocean



When a packet arrives at a router



Routers use a protocol called **BGP** to decide what router the packet should go to next:



The route it takes to get from A → B might be different from B → A.

Exercise: Run `traceroute google.com` to see what steps your packet takes to get to google.com.

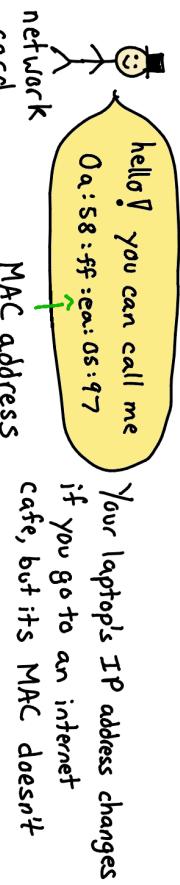
Local networking

how to talk to a computer in the same room

Every computer is in a subnet. Your subnet is the list of computers you can talk to directly.



What does it mean to talk "directly" to another computer? Well, every computer on the internet has a network card with a MAC address.



When you send a packet to a computer in your subnet, you put the computer's MAC address on it. To

get the right MAC, your

computer uses a protocol called ARP: (Address Resolution Protocol)

I will store that in my ARP table!

You can run `arp -na` to see the contents of the

ARP table on your computer. It should look like this:

```
$ arp -na
(192.168.1.120) at 94:53:30:91:98:c8 [ether] on wlp3s0
  m7: wifi
  ? (MAC for 192.168.1.120 (my printer))
```

There are 2 kinds of DNS servers:

recursive

I can get you an IP address for ANY website by asking the right authoritative DNS server (like art.ns.cloudflare.com)

wanna know where jvns.ca is?

Talk to ME! DNS server (like art.ns.cloudflare.com)

When you query a recursive DNS server, here's what happens:

I have to talk to THREE authoritative DNS servers?

Okay!

Where's jvns.ca?

ask there!

where's jvns.ca?

root DNS server
 a.root-servers.net

.ca DNS server
 jvns.ca
 jvns.ca-servers.ca

art.ns.cloudflare.com
 jvns.ca
 art.ns.cloudflare.com

Recursive DNS servers usually cache DNS records.

Every DNS record has a TTL ("time to live") that says how long to cache it for. You often can't force them to update their cache. You just have to wait:

I updated my DNS records, but when I visit the site in my browser I see the old version :)

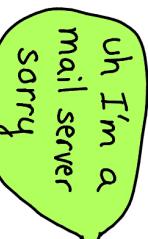
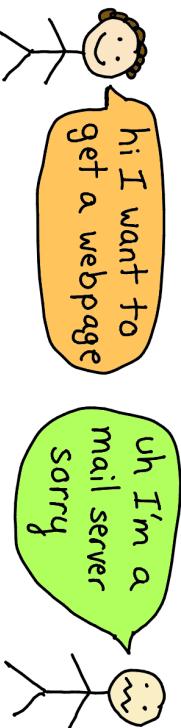
20 minutes later after the recursive DNS server cache updates...

everything is great now

What's a "port"?

ports are part of the **TCP** and **UDP** protocols
(**TCP** port 999 and **UDP** port 999 are different)

When you send a **TCP** message, you want to talk to a specific kind of program
This would be bad:



We want to have different kinds of programs on the same server:

minecraft **DNS** **email**

So every **TCP** packet has a part number between 1 and 65535 on it:

I'm listening on **TCP** port 80

here's a **TCP** packet with part 80 on it!

ooh! that's for ME!

netstat and lsof can tell you which ports are in use on your computer

Some common ports:
HTTP: **TCP** port 80
HTTPS: **TCP** port 443
SMTP: **TCP** port 25
(mail)
Minecraft: **TCP + UDP** 25565

Sockets

Step ②: now that we have an IP address, the next step is to open a socket!

Let's learn what that is.

your program doesn't know how to do **TCP**

idk what "**TCP**" is I just want to get a webpage

OS
code.py
don't worry!
I can help!

to use **TCP**
TCP
raw
for ULTIMATE POWER, ping uses this to send ICMP packets

4 common socket types

TCP
to use **UDP**
UDP

UNIX
to talk to same computer

what using sockets is like
step 1: ask the OS for a socket
step 2: connect the socket to an IP address and port
step 3: write to the socket to send data

when you connect with a **TCP** socket



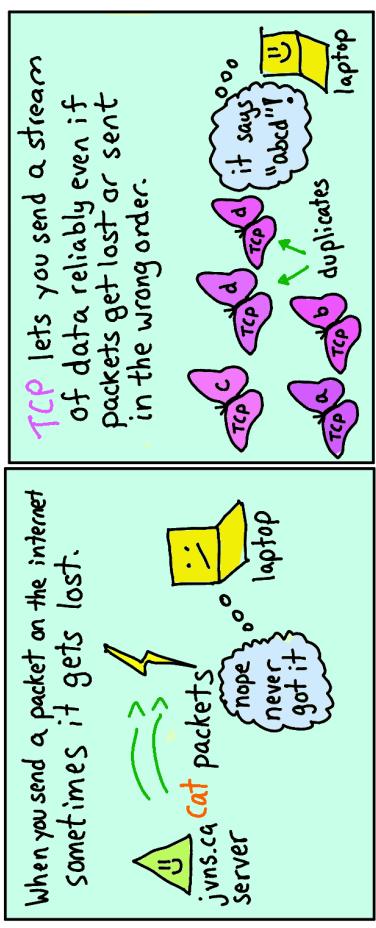
(we'll explain this **SYN ACK** thing soon)

When you write to a socket
code.py → writes lots of data ♥♥♥♥
→ splits it up into packets → into packets to send it

this socket interface is great! the operating system does so much for me!

TCP: how to reliably get a cat

Step ③ in our plan is "open a TCP connection!"
Let's learn what this "TCP" thing even is!)



how does TCP work, you ask? WELL!

how to know what order the packets should go in:

Every packet says what range of bytes it has

Like this:

Once upon a time there was a magical oyster
The position of the first byte (0, 14, 30 in our example) is called the "sequence number"

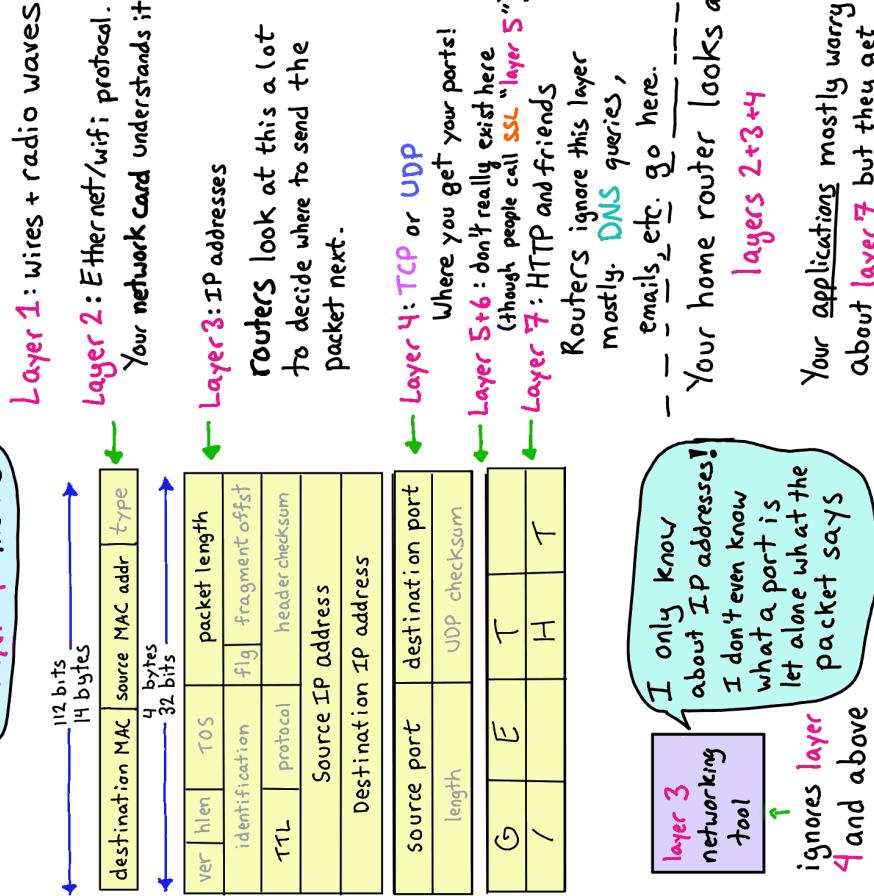
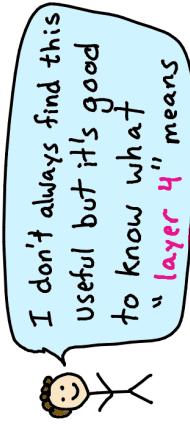
Then the client can assemble all the pieces into:

"Once upon a time there was a magical oyster"

If the server doesn't get an ACKnowledgement, it will retry sending the data.

networking layers

Networking layers mostly correspond to different sections of a packet.



The cool thing is that the layers are mostly independent of each other - You can change the IP address (layer 3) and not worry about layers 4+7

Your applications mostly worry about layer 7 but they get to tell the operating system what IP and port to use.

The network card in your computer only cares about layers 1+2.



... and now for even MORE



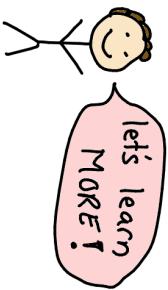
This is what a **TCP** header looks like:

Source Port	Destination Port
Sequence Number	
Acknowledgement Number	
Offset	Reserve
Data	RCV-SQNS
Checksum	ACK-NSE
Options	Window
Urgent Pointer	PSH-URG
Padding	

We've covered the basics of how to download a cat picture now! But there's a lot more to know! Let's talk about a few more topics.

We'll explain a little more about networking protocols:

- what a port actually is
- how a packet is put together
- security: how **SSL** works
- the different networking layers
- **UDP** and why it's amazing
- and how packets get sent from place to place:
- how packets get sent in a local network
- and how packets get from your house to jvns.ca
- networking notation



When you see "connection refused" or "connection timeout" errors, that means the **TCP** handshake didn't finish!

I ran `sudo tcpdump host jvns.ca` in one terminal and `curl jvns.ca` in another. This is some of the output:

```
localhost:51104 > 104.28.6.94:80 Flags [S] 104.28.6.94:80 > localhost:51104 Flags [S, ACK] localhost:51104 > 104.28.6.94:80 Flags [.]
```

jvns.ca IP address
• S is for SYN
• is for ACK



But what do "SYN" and "ACK" mean? Well! **TCP** headers have 6 bit flags (**SYN**, **ACK**, **RST**, **FIN**, **PSH**, **URG**) that you can set (you can see them in the diagram.) A **SYN** packet is a packet with the **SYN** flag set to 1.

When you see "connection refused" or "connection timeout" errors, that means the **TCP** handshake didn't finish!

"The TCP Handshake"

HTTP

important HTTP headers

Step ④: Finally, we can request `cat.png`!

Every time you get a webpage or see an image online, you're using HTTP:

HTTP is a pretty simple plaintext protocol. In fact, it's so simple that you can make a HTTP request by hand right now. Let's do it!!!

First, let's make a file called `request.txt`:

```
GET / HTTP/1.1
Host: ask.metafilter.com
User-Agent: zine
(put 2 newlines at the end)
```

Then:

```
cat request.txt | nc metafilter.com 80
```

The `nc` command ("net `cat`") sets up a `TCP` connection to `metafilter.com` and sends the HTTP request you wrote! The response we get back looks like:

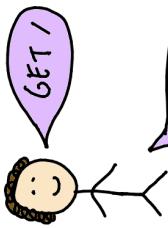
```
200 OK
Content-Length: 120321
...headers...
a bunch of
HTML
```

HTTP/2 is the next version of HTTP. It's very different but we're out of space.

This is a HTTP request:
`GET /cat.png HTTP/1.1`
`Host: jvns.ca`
`User-Agent: zine`

The User-Agent and Host lines are called "headers". They give the webserver extra information about what webpage you want!

the Host header - my favorite!



dude, do you even know
how many websites I
serve? You gotta be
more specific.

GET /
Host: jvns.ca

Now we're talking

Most servers serve lots of
different websites. The
Host header lets you pick
the one you want!

More useful headers:

Cookie

When you're logged
into a website,
your browser sends data
in this header! This
is how the server knows
you're logged in.

Accept-Encoding

Want to save
bandwidth? Set
this to "gzip" and
the server might
compress your
response.

User-Agent

Lots of servers use
this to check if you're
using an old browser
or if you're a bot.